



Iterative Soft Thresholding for LASSO (ISTA)

Oscar Ma, Max Maejima, Alan Zhong



What is LASSO?

- LASSO stands for Least Absolute Shrinkage and Selection Operator
- Modified version of the Least Squares Regression Model
 - LASSO is used for regression problems similar to ordinary least squares but uses additional parameters
 - Lambda is the regularization parameter
 - L1 Norm of Beta is the sum of the absolute value of the Beta coefficients

Why do we use LASSO?

- LASSO is a modified version of the Least Squares Regression Model
- Linear Regression with $\{X_1, \dots, X_n\}$ where n is large often overfits the data
- In real world examples, some X_i 's are often statistically irrelevant
- LASSO's optimization problem emphasizes the most important independent variables in our analysis, attempting to avoid large coefficients when possible

LASSO Optimization Formula

$$\hat{\beta} = \arg \min_{\beta} \{ \|y - X\beta\|_2^2 + \lambda \|\beta\|_1 \}$$

- **Similar to Multiple Linear Regression Optimization,**

Our Beta Vector is the coefficients $\{\beta_0, \beta_1, \dots, \beta_i\}$ for our variables $\{X_1, \dots, X_i\}$

β_0 is the intercept (It is the predicted value when all $X_i = 0$ for all i)

- Our y-vector contains the true y-values for the variable we are aiming to predict
- X is the matrix that contains the values of the X_i 's that coincide with our actual Y_i values as rows
- **Unique to LASSO**, Lambda is a non-negative real number that controls the strength of LASSO's influence on our optimization problem.
 - When **Lambda = 0**, this is an ordinary Least Squares Regression problem
 - When **Lambda is large**, many coefficients in our Beta Vector are forced to 0, extracting the strongest of X_i predictor variables for regression

The Beta Vector

$$y_i = \beta_0 + \sum_{j=1}^p \beta_j x_{ij}$$

- Elements of beta are coefficients of a linear polynomial to predict elements in vector y
 - Each element of the Beta Vector β_j corresponds to predictor variable X_j
-
- One way to initialize β is the Zero Vector
 - Treating all coefficients equally
 - LASSO's goal is to drive some coefficients to 0 for variable selection; it makes sense to start iterating with the zero vector

Importance of the L1-Norm

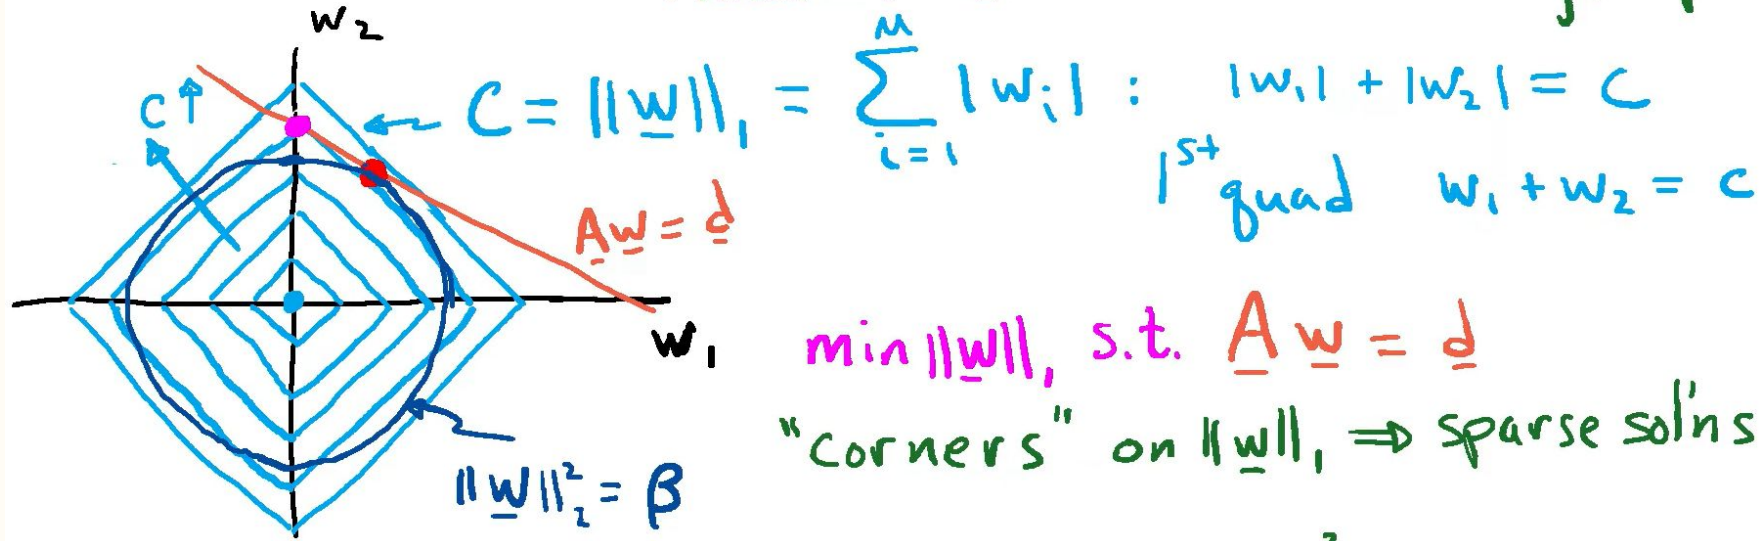
$$\|\beta\|_1 = \sum_{j=1}^p |\beta_j|$$

- L1-Norm of Beta - the **sum** of all the absolute coefficients of the X_i 's
- This optimization problem requires **minimization**
 - Adding a factor of the L1-Norm to the objective function can drastically change the optimized Beta Vector
 - Compared to a Least Squares Minimization problem, LASSO prefers pushing β_j to 0 for less relevant variables X_j
 - This is known as the L1 Penalty
- But we have a PROBLEM!
 - The L1 Norm is not differentiable at 0 due to Absolute Values

Convex relaxation gives tractable problem 3

$\min_{\underline{w}} \|\underline{w}\|_1$, s.t. $\|\underline{A}\underline{w} - \underline{d}\|_2^2 < \varepsilon$ LASSO: Least

Absolute Selection & Shrinkage Operator



$\min \|\underline{w}\|_1$, s.t. $\underline{A}\underline{w} = \underline{d}$

"corners" on $\|\underline{w}\|_1 \Rightarrow$ sparse sol's

$\min \|\underline{w}\|_2^2$ s.t. $\underline{A}\underline{w} = \underline{d}$ circular $\|\underline{w}\|_2^2 \Rightarrow$ non sparse solutions

Iterative Soft Thresholding for LASSO


- To solve this L1-Norm differentiation issue, we use **Iterative Soft Thresholding**

Step 1: Gradient Descent

- Take the Gradient of the smooth Least Squares part of our objective function

Step 2: Soft Thresholding

- Shrinks smaller coefficient values toward 0
- Handle Non-Differentiability of the L1 Penalty



Gradient Descent

- First, split the objective function into two parts
 - A Smooth part, denoted $f(\beta)$, which is our OLS objective function
 - A Non-Smooth part, denoted $g(\beta)$

$$f(\beta) = \|y - X\beta\|_2^2.$$

$$g(\beta) = \lambda \|\beta\|_1$$

- Then, find the gradient of f
 - Expanded, $f(\beta) = (y - X\beta)^T(y - X\beta) = y^T y - 2\beta^T X^T y + \beta^T X^T X \beta$
 - Taking the derivative with respect to β :
 - $d/d\beta (y^T y) = 0$
 - $d/d\beta (2\beta^T X^T y) = 2X^T y$
 - $d/d\beta (\beta^T X^T X \beta) = 2X^T X \beta$

$$\nabla f(\beta) = 0 - 2X^T y + 2X^T X \beta = 2X^T (X\beta - y).$$

Intermediate Step for Proximal Operator

$$\tilde{\beta}^{(k)} = \beta^{(k)} - \alpha \nabla f(\beta^{(k)}).$$

- $\beta^k \rightarrow \tilde{\beta}^k \rightarrow \beta^{k+1}$
 - Big Idea: Take a step in the direction of the gradient
 - High number of iterations \rightarrow use a small fixed alpha (we used 10^{-3})
- This is our **input vector** for the proximal operator

Proximal Operator Function

$$\beta^{(k+1)} = \text{prox}_{\alpha_k g} \left(\beta^{(k)} - \alpha_k \nabla f(\beta^{(k)}) \right).$$

- For LASSO, we choose Soft Thresholding as our Proximal Operator, hence **Iterative Soft Thresholding Algorithm (ISTA)**
- Soft Thresholding Algorithm

$$\beta^{(k+1)} = \text{soft}(\beta^{(k)} - \alpha \nabla f(\beta^{(k)}), \alpha \lambda),$$

Where each element of

$$\beta_j^{(k+1)} = \text{sign}(\tilde{\beta}_j^{(k)}) \max(|\tilde{\beta}_j^{(k)}| - \alpha \lambda, 0).$$

Soft Thresholding

$$\beta_j^{(k+1)} = \text{sign}(\tilde{\beta}_j^{(k)}) \max\left(|\tilde{\beta}_j^{(k)}| - \alpha \lambda, 0\right).$$

- The ISTA Algorithm:
 - Input: jth element of intermediate step β_j^k -tilde
 - Output: jth element of β^{k+1}

$$\beta_j^k\text{-tilde} < -\alpha\lambda \text{ ---> } \beta_j^{k+1} := \beta_j^k\text{-tilde} + \alpha\lambda$$

$$-\alpha\lambda < \beta_j^k\text{-tilde} < \alpha\lambda \text{ ---> } \beta_j^{k+1} := 0$$

$$\beta_j^k\text{-tilde} > \alpha\lambda \text{ ---> } \beta_j^{k+1} := \beta_j^k\text{-tilde} - \alpha\lambda$$

$$\min_{\beta} \|y - X\beta\|_2^2 + \lambda \|\beta\|_1, \beta \in \mathbb{R}^d$$

β = Coeffs X = Predictor Variable Matrix y = Predicted variable vector

Let $f(\beta) = \|y - X\beta\|_2^2$, $g(\beta) = \lambda \|\beta\|_1$ where $\|\beta\|_1 = \sum_{i=1}^d |\beta_i|$

$$\nabla f(\beta) = -2X^T(y - X\beta)$$

$$\tilde{\beta} = \beta - \alpha \nabla f(\beta) \quad \alpha = \text{Step Size}$$

$$\beta' = \text{prox}_{\alpha} g(\tilde{\beta})$$


$$\beta'_j = \text{sign}(\tilde{\beta}_j) \cdot \max(|\tilde{\beta}_j| - \alpha\lambda, 0)$$


where if $|\tilde{\beta}_j| < \alpha\lambda \Rightarrow \beta'_j = 0$

Let $\beta^0 = \vec{0}$

$$\nabla f(\beta^0) = -2X^T y \Rightarrow \tilde{\beta}^0 = -\alpha 2X^T y$$

$$\beta'_j = \text{sign}(\tilde{\beta}^0_j) \cdot \max(|\tilde{\beta}^0_j| - \alpha\lambda, 0)$$





Implementation

We applied Lasso Regression using the Iterative Soft Thresholding Algorithm (ISTA) to identify key factors influencing career outcomes based on a dataset of student profiles.

```
def soft_thresholding(x, alpha, lambda_):  
    """Applies soft-thresholding to shrink values toward  
    zero."""  
    return np.sign(x) * np.maximum(np.abs(x) -  
alpha*lambda_, 0)
```

Implementation Continued

```
def ista(X, y, lambda_, alpha=1e-3, max_iter=1000, tol=1e-6):  
    m, n = X.shape  
    beta = np.zeros(n)    # Initialize beta coefficients to zero  
  
    for _ in range(max_iter):  
        gradient = X.T @ (X @ beta - y) / m    # Compute the gradient of loss  
        beta_new = soft_thresholding(beta - alpha * gradient, alpha, lambda_)  
  
        if np.linalg.norm(beta_new - beta, ord= 2) < tol:  
            break    # Check if the update is small enough to consider convergence  
  
        beta = beta_new  
  
    return beta
```


Issue with Cross Validation for Choosing λ

A common way to choose parameters in machine learning is cross validation, a method of iterating over possible parameter values and selecting one that minimizes the squared error.

With LASSO we are balancing two priorities; a Beta with sparse weights and reducing squared error. We found using cross validation results in bias towards smallest λ possible. Thus in LASSO it is important to recognize minimizing error alone doesn't guarantee a sparse model and chosen λ may lead to overfitting or dense models, defeating the purpose of feature selection.

$$\hat{\beta} = \arg \min_{\beta} \{ \|y - X\beta\|_2^2 + \lambda \|\beta\|_1 \}$$

Data

Our dataset includes academic, personal, and career information for recent graduates:



Results

Feature Importance (ISTA Lasso) for Job_Offers
(lambda = 0.01):

Projects_Completed	0.012913
Age	0.009617
Internships_Completed	0.008134
SAT_Score	0.004453
University_GPA	0.003893
Gender_Other	0.003813
University_Ranking	0.003028
High_School_GPA	0.002633
Soft_Skills_Score	0.002628

Networking_Score
0.001614

Field_of_Study_Engineering
0.001511

Field_of_Study_Business
0.001395

Certifications 0.000381

Gender_Male 0.000000

**Field_of_Study_Computer Science
0.000000**

**Field_of_Study_Law
0.000000**

**Field_of_Study_Mathematics
0.000000**

**Field_of_Study_Medicine
0.000000**



Results

Feature Importance (ISTA Lasso) for Starting_Salary
(lambda = 0.05):

Internships_Completed	0.013612
Field_of_Study_Law	0.009995
University_Ranking	0.009057
Certifications	0.008917
Field_of_Study_Medicine	0.008152
Field_of_Study_Business	0.003515
Soft_Skills_Score	0.002962
Gender_Male	0.002493
High_School_GPA	0.001999

University_GPA	0.001986
Projects_Completed	0.000467
Field_of_Study_Computer Science	0.000261
Age	0.000000
SAT_Score	0.000000
Networking_Score	0.000000
Gender_Other	0.000000
Field_of_Study_Engineering	0.000000
Field_of_Study_Mathematics	0.000000



References

1. Barry Van Veen. (n.d.). Solving l_1 Regularized Least Squares via Proximal Gradient Descent. YouTube.
<https://www.youtube.com/watch?v=CBMDwMESuWs&list=PLGI7M8vwfrFMo6i-kmYROMN-GUdKNHGLv&index=4>
 2. Murphy, K. P. (2012). *Machine learning : a probabilistic perspective*. Cambridge, MA: MIT. ISBN: 9780262018029 0262018020
 3. Professor Brian Powers. (4/10/2024).
UW Madison Department of Statistics. brpowers2@wisc.edu
 4. ChatGPT
- 